


Credit Card fraud detection using Isolation Forest

Spoorthy Reddy Jarugu

Vellore Institute of Technology
India

September 29, 2023





Overview

- [Problem Definition](#)
- [Data Preprocessing and Visualization](#)
- [Model built and Prediction](#)
- [Conclusion](#)

Problem Definition

Credit card fraud detection

Dataset Description

Data Preprocessing and Visualization

Loading the data

Cleaning the data

Statistical measures of the feature Class

Correlation

Model built and Prediction

Isolation Forest

Local Outlier Factor

Classification Report

Analysis

Evaluation Result

Conclusion



TULIP

Team for Universal Learning and Intelligent Processing



Problem Definition

Credit card fraud detection

Dataset Description

Data Preprocessing and Visualization

Model built and Prediction

Conclusion

Problem Definition



Credit card fraud detection

Problem Definition

Credit card fraud detection

Dataset Description

Data Preprocessing and Visualization

Model built and Prediction

Conclusion

Defn

Credit Cards are the most commonly used mode of payment nowadays. The reason is it has multiple features, which make it easy for users to make payments on the spot. Credit card fraud can be defined as any unauthorized use of a credit card, such as using a stolen credit card or making unauthorized purchases with a valid credit card. This dataset is taken from Kaggle.

- The dataset contains transactions made by credit cards in **September 2013** by European cardholders.
- Unfortunately, due to confidentiality issues, original features are not given. The features provided are the result of the PCA transformation.
- There are a variety of techniques that can be used to detect credit card fraud. One common technique is to use machine learning models to identify patterns in fraudulent transactions.



Dataset Description

- Problem Definition
- Credit card fraud detection
- Dataset Description**
- Data Preprocessing and Visualization
- Model built and Prediction
- Conclusion

Name	Count
Rows	284807
Columns	31

Rows

- This dataset contains 284807 rows of vales.
- The entire dataset has 284315 Valid transactions, and 492 are Fraud transactions.

Columns

- The total number of columns present in the dataset is 31
- Time - Time of the transaction happened, V1 to v28 - are the principle component obtained with PCA, Amount - transaction amount, Class - valid or fraud transaction



[Problem Definition](#)

[Data Preprocessing and Visualization](#)

[Loading the data](#)

[Cleaning the data](#)

[Statistical measures of the feature](#)

[Class](#)

[Correlation](#)

[Model built and Prediction](#)

[Conclusion](#)

Data Preprocessing and Visualization



TULIP

Team for Universal Learning and Intelligent Processing



Loading the data

- [Problem Definition](#)
- [Data Preprocessing and Visualization](#)
- [Loading the data](#)**
- [Cleaning the data](#)
- [Statistical measures of the feature](#)
- [Class](#)
- [Correlation](#)
- [Model built and Prediction](#)
- [Conclusion](#)

- After the import statements, the initial step is to load the data
- Setting the proper path where the data is available is an important step

```
data = pd.read_csv('./dataset/creditcard.csv')  
data.head()
```

Figure 1: Loading the data





Cleaning the data

- Problem Definition
- Data Preprocessing and Visualization
- Loading the data
- Cleaning the data
- Statistical measures of the feature Class
- Correlation
- Model built and Prediction
- Conclusion

- Checking for missing values of each column
- Checking for the data type of each column

```
data.isnull().sum()
```

Time	0
V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0

Figure 2: Checking for missing values

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 284807 entries, 0 to 284806  
Data columns (total 32 columns):  
#   Column  Non-Null Count  Dtype  
---  -  
0   Time    284807 non-null float64  
1   V1       284807 non-null float64  
2   V2       284807 non-null float64  
3   V3       284807 non-null float64  
4   V4       284807 non-null float64  
5   V5       284807 non-null float64  
6   V6       284807 non-null float64  
7   V7       284807 non-null float64  
8   V8       284807 non-null float64  
9   V9       284807 non-null float64  
10  V10      284807 non-null float64  
11  V11      284807 non-null float64  
12  V12      284807 non-null float64  
13  V13      284807 non-null float64  
14  V14      284807 non-null float64  
15  V15      284807 non-null float64  
16  V16      284807 non-null float64  
17  V17      284807 non-null float64  
18  V18      284807 non-null float64  
19  V19      284807 non-null float64
```

Figure 3: Data type of each column



Visual representation of the feature class

- [Problem Definition](#)
- [Data Preprocessing and Visualization](#)
- [Loading the data](#)
- [Cleaning the data](#)**
- [Statistical measures of the feature Class](#)
- [Correlation](#)
- [Model built and Prediction](#)
- [Conclusion](#)

- Number of classes with respect to frequency that are Valid transaction and Fraud transaction




Figure 4: Class Visualization

- From the above diagram, we can see that Valid transactions are way more significant than fraud transactions



Statistical measures of the feature Class

- Problem Definition
- Data Preprocessing and Visualization
- Loading the data
- Cleaning the data
- Statistical measures of the feature Class
- Correlation
- Model built and Prediction
- Conclusion

- Diving the feature 'Class==0' as a Valid dataset.

```
print("Details of Valid transaction")
Valid.Amount.describe()
```

Details of Valid transaction

count	284315.000000
mean	88.291022
std	250.105092
min	0.000000
25%	5.650000
50%	22.000000
75%	77.050000
max	25691.160000
Name: Amount, dtype: float64	

Figure 5: Details of Valid transaction

- Diving the feature 'Class==1' as a Fraud dataset.

```
print("Details of the Fraud transaction")
Fraud.Amount.describe()
```

Details of the Fraud transaction

count	492.000000
mean	122.211321
std	256.683288
min	0.000000
25%	1.000000
50%	9.250000
75%	105.890000
max	2125.870000
Name: Amount, dtype: float64	

Figure 6: Details of Fraud transaction



Valid and Fraud transactions with respect to Amount

- Problem Definition
- Data Preprocessing and Visualization
 - Loading the data
 - Cleaning the data
 - Statistical measures of the feature Class
 - Correlation
- Model built and Prediction
- Conclusion

- Trying to visually see how many different transactions are for Fraud and Valid in terms of Amount.
- From the visual representation, we can see Fraud transactions are of a small amount compared to Valid transactions.




Figure 7: Transactions with respect to Amount



Valid and Fraud transactions with respect to Time

- Problem Definition
- Data Preprocessing and Visualization
 - Loading the data
 - Cleaning the data
 - Statistical measures of the feature Class
 - Correlation
- Model built and Prediction
- Conclusion

- Trying to visually see how many different transactions are for Fraud and Valid in terms of Time.




Figure 8: Transactions with respect to Time



Correlation

[Problem Definition](#)

[Data Preprocessing and Visualization](#)

[Loading the data](#)

[Cleaning the data](#)

[Statistical measures of the feature Class](#)

[Correlation](#)

[Model built and Prediction](#)

[Conclusion](#)

- Correlation is a statistical measure used to determine if there is a relationship between two variables and how strongly that is related.
- Correlation coefficients range from -1 to +1. A correlation coefficient of -1 indicates a perfect negative correlation, which means that the two variables move in opposite directions. A correlation coefficient of +1 indicates a perfect positive correlation, which means that the two variables move in the same direction.
- A correlation coefficient of 0 indicates no correlation, meaning there is no relationship between the two variables.
- To avoid multicollinearity and improve model performance, we will remove the highly correlated variables to reduce the redundancy of the two variables.



TULIP

Team for Universal Learning and Intelligent Processing



Correlation

Problem Definition

Data Preprocessing and Visualization

Loading the data

Cleaning the data

Statistical measures of the feature

Class

Correlation

Model built and Prediction

Conclusion

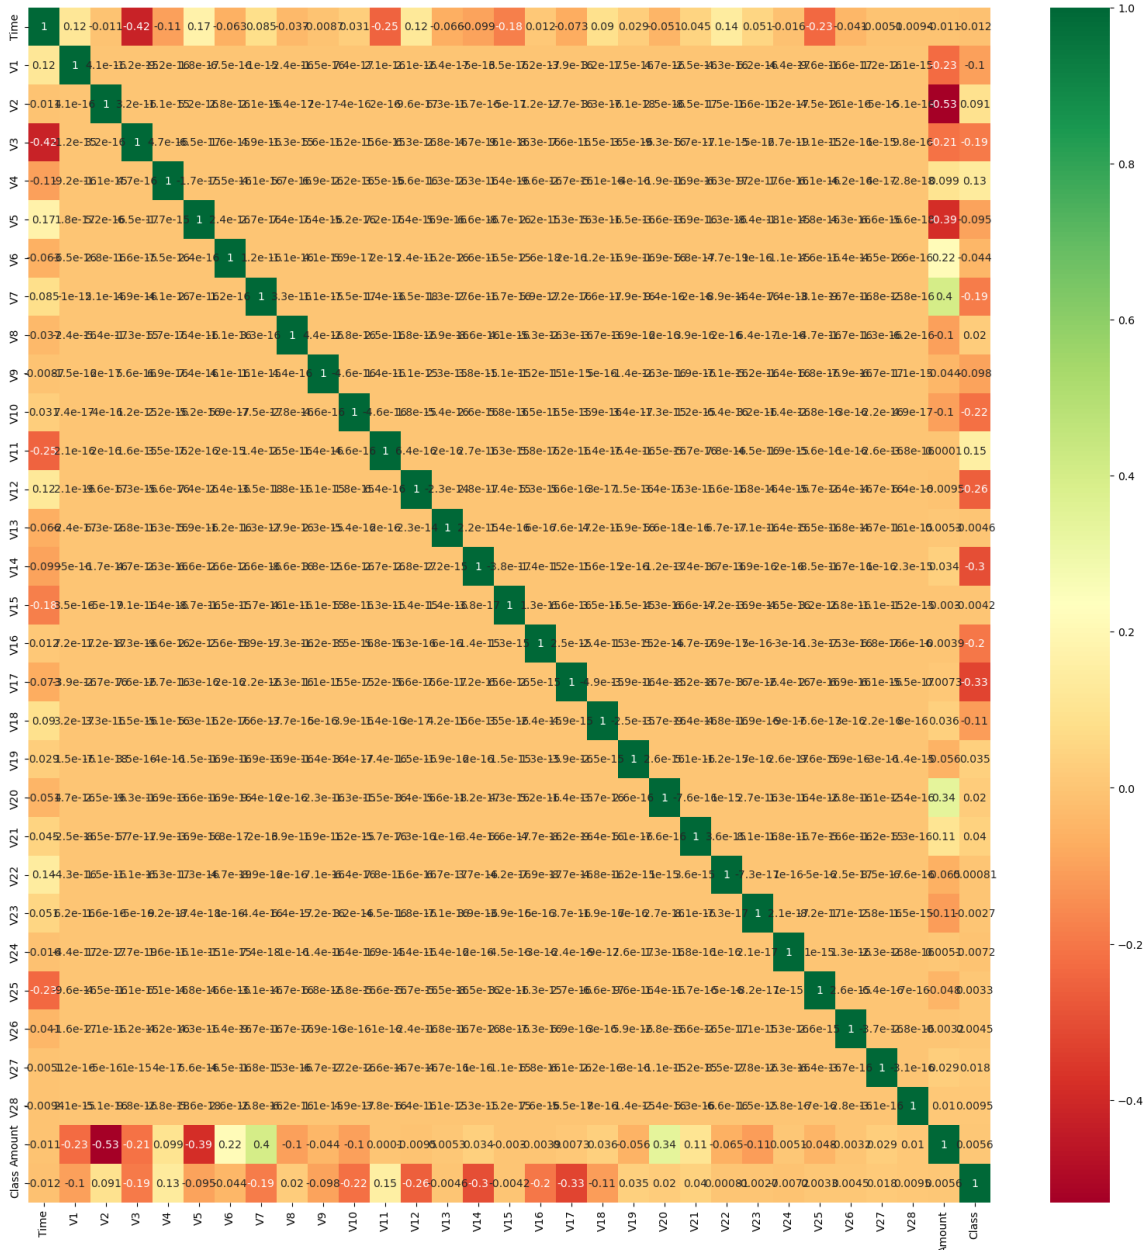


Figure 9: Correlation matrices



- [Problem Definition](#)
- [Data Preprocessing and Visualization](#)
- [Model built and Prediction](#)**
- [Isolation Forest](#)
- [Local Outlier Factor](#)
- [Classification Report](#)
- [Analysis](#)
- [Evaluation Result](#)
- [Conclusion](#)

Model built and Prediction



Isolation Forest

- Problem Definition
- Data Preprocessing and Visualization
- Model built and Prediction
- Isolation Forest**
- Local Outlier Factor
- Classification Report
- Analysis
- Evaluation Result
- Conclusion

- The Isolation Forest algorithm is a powerful tool for anomaly detection. It is fast, efficient, and robust to outliers and noise. It can be used to detect anomalies in various applications, including fraud detection, intrusion detection, medical diagnosis, network monitoring, and financial market analysis.
- The Isolation Forest algorithm is an unsupervised anomaly detection algorithm that works by isolating anomalies by randomly selecting features and split values. The algorithm builds a forest of isolation trees, where each tree is trained on a random sample of the data. The algorithm then calculates the anomaly score for each data point by measuring how deep it is in the forest. Anomalies are typically assigned lower anomaly scores than normal data points.
- Here, for this task, we are using the Isolation Forest for Credit card fraud detection.



Isolation Forest

Problem Definition
Data Preprocessing and Visualization
Model built and Prediction
Isolation Forest
Local Outlier Factor
Classification Report
Analysis
Evaluation Result
Conclusion

- This is how we create new instances for isolation forest
`IsolationForest(n_estimators=1000, max_samples="auto", contamination=0.1, random_state=42)`
 - ◆ `n_estimators`: This parameter specifies the number of trees to build in the forest. A higher value of `n_estimators` will result in a more robust model but will also take longer to train.
 - ◆ `max_samples`: This parameter specifies the maximum number of samples to be used for training each tree in the forest. An "auto" value means that all the samples will be used to train each tree.
 - ◆ `contamination`: This parameter specifies the proportion of outliers that the IsolationForest model is expected to find in the data. A value of 0.1 means that the model is expected to find 10% of the data points to be outliers.
 - ◆ `random_state`: This parameter specifies the random seed to use. Setting the random state to a fixed value will ensure the model produces consistent results each time it is trained.



Local Outlier Factor

Problem Definition
Data Preprocessing and Visualization
Model built and Prediction
Isolation Forest
Local Outlier Factor
Classification Report
Analysis
Evaluation Result
Conclusion

- Local Outlier Factor (LOF) is an unsupervised anomaly detection algorithm that identifies outliers based on their local density. LOF is calculated by comparing the local density of a data point to the local densities of its neighbours.
- In LOF a data point is considered to be an outlier if it has a significantly lower local density than its neighbors.
- The new instance of Local Outlier Factor is created as follows:
`LocalOutlierFactor(n_neighbors=50,leaf_size=10, contamination=0.1)`
 - ◆ `n_neighbors`: This parameter specifies the number of nearest neighbours to consider when calculating the local density of a data point. A higher value of `n_neighbors` will result in smoother local density estimates.
 - ◆ `leaf_size`: This parameter specifies the maximum number of data points that can be stored in a single node of the tree used to calculate the nearest neighbours of a data point. A higher value of `leaf_size` will result in faster computation, but it may also lead to less accurate results.
 - ◆ `contamination`: This parameter specifies the proportion of outliers the `LocalOutlierFactor` model is expected to find in the data.



Comparing IF and LOR

- Problem Definition
- Data Preprocessing and Visualization
- Model built and Prediction
 - Isolation Forest
 - Local Outlier Factor
- Classification Report
- Analysis
- Evaluation Result
- Conclusion

- Both the algorithms gave a reasonable accuracy rate. However, Isolation Forest (IF) effectively fetches outliers for large datasets, whereas Local Outlier Factor (LOR) algorithms are computationally expensive for large datasets.
- IF is generally more interpretable than LOF.
- LOF has more hyperparameters to tune than IF.

Algorithm	Accuracy
Isolation Forest	0.997156
Local Outlier Factor	0.996524

Figure 10: Accuracy Report



Classification Report

- Problem Definition
- Data Preprocessing and Visualization
- Model built and Prediction
- Isolation Forest
- Local Outlier Factor
- Classification Report
- Analysis
- Evaluation Result
- Conclusion

- Classification report is the performance of an Isolation Forest model on a given dataset.
 - ◆ Precision is the fraction of positive predictions that are positive.
 - ◆ Recall is the fraction of actual positives that are correctly identified.
 - ◆ F1 score is a harmonic mean of precision and recall.
 - ◆ Support is the total number of examples in a given class.

Classification Report :					
	precision	recall	f1-score	support	
0	1.00	1.00	1.00	28432	
1	0.26	0.27	0.26	49	
accuracy				1.00	28481
macro avg				0.63	28481
weighted avg				1.00	28481

Figure 11: Classification Report



Classification Report

Problem Definition
Data Preprocessing and Visualization
Model built and Prediction
Isolation Forest
Local Outlier Factor
Classification Report
Analysis
Evaluation Result
Conclusion

- The model has perfect accuracy on the Valid Transaction (i.e., class 0), with a precision of 1.0 and a recall of 1.0. However, the model performs poorly on the Fraud transaction (class 1), with a precision of 0.26 and a recall of 0.27.
- Overall, the classification report shows that the model has good performance on the majority class that is a Valid transaction, but poor performance on the minority class that is a Fraud transaction. This is a common problem with classification models, which are often trained on datasets with imbalanced class distributions.
- Here are some ways to improve the performance of the model on the minority class:
 - ◆ Use a weighted loss function to give more weight to the minority class.
 - ◆ Use oversampling or undersampling techniques to balance the class distribution of the training set.
 - ◆ Use a different classification algorithm better suited for imbalanced datasets.



Analysis

Problem Definition
Data Preprocessing and Visualization
Model built and Prediction
Isolation Forest
Local Outlier Factor
Classification Report
Analysis
Evaluation Result
Conclusion

- From the above visual representation of the class feature, we can understand that Valid transactions are more in number than Fraud transactions. When understanding the percentage, it seems 0.17% are fraud transactions in the entire data.
- As it is very low, this data is highly imbalanced data.
- In general, to balance the dataset, we use two methods:
 - ◆ Undersampling: This approach involves reducing the number of samples in the majority class by randomly removing samples. This can be done until the majority class has the same number of samples as the minority class.
 - ◆ Oversampling: This approach involves increasing the number of samples in the minority class by creating synthetic samples. This can be done using various techniques, such as SMOTE (Synthetic Minority Over-sampling Technique).



Evaluation Result

- Problem Definition
- Data Preprocessing and Visualization
- Model built and Prediction
 - Isolation Forest
 - Local Outlier Factor
 - Classification Report
 - Analysis
 - Evaluation Result**
- Conclusion

- Undersampling the dataset is done to get accurate results for IF and LOF algorithms.
- We can see that the IF algorithm outperformed the LOF algorithm and obtained an accuracy of 97.5%. Whereas LOF obtained 87.8%.
- Overall, the classification report shows that the model performed very well on class 0 data and moderately well on class 1 data. The model has a high accuracy, precision, and recall for class 0 data. The model has a moderate accuracy, precision, and recall for class 1 data.

Algorithm	Accuracy
Isolation Forest	0.9756
Local Outlier Factor	0.878647

Figure 12: Accuracy Report

- Outliers obtained are 350

Isolation Forest: 350				
Classification Report :				
	precision	recall	f1-score	support
0	0.99	0.99	0.99	15000
1	0.64	0.65	0.65	492
accuracy			0.98	15492
macro avg	0.82	0.82	0.82	15492
weighted avg	0.98	0.98	0.98	15492

Figure 13: Classification Report



[Problem Definition](#)

[Data Preprocessing and Visualization](#)

[Model built and Prediction](#)

Conclusion

Conclusion



Conclusion

[Problem Definition](#)

[Data Preprocessing and Visualization](#)

[Model built and Prediction](#)

[Conclusion](#)

- Flip 01 task I have chosen is to find "Credit Card Fraud Detection using Isolation Forest"
- The Isolation Forest algorithm has a number of advantages over other anomaly detection algorithms :
 - ◆ It is robust to outliers and noise.
 - ◆ It can be used to detect anomalies in both high-dimensional and low-dimensional data.
 - ◆ It is easy to implement and interpret.
- Using Isolation Forest for detecting anomalies gives very good accuracy compared to other algorithms.



TULIP

Team for Universal Learning and Intelligent Processing



Questions?

- [Problem Definition](#)
- [Data Preprocessing and Visualization](#)
- [Model built and Prediction](#)
- [Conclusion](#)



Contact Information

Spoorthy Reddy Jarugu
Vellore Institute of Technology
Vellore, India

